


```
In [2]: import keras
from keras.models import Sequential # Import Sequential from keras.models
from keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense
from keras.datasets import mnist
%matplotlib inline
```


```
In [3]: #Load mnist dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
 11490434/11490434 [=====] - 4s 0us/step

```
In [4]: img_rows, img_cols = 28, 28
```

```
In [5]: import matplotlib.pyplot as plt
fig = plt.figure()
for i in range(9):
    plt.subplot(3,3,i+1)
    plt.tight_layout()
    plt.imshow(X_train[i], cmap='gray', interpolation='none')
    plt.title("Digit: {}".format(y_train[i]))
    plt.xticks([])
    plt.yticks([])
fig
```





```
In [6]: if keras.backend.image_data_format() == 'channels_first':
        X_train = X_train.reshape(X_train.shape[0], 1, img_rows, img_cols)
        X_test = X_test.reshape(X_test.shape[0], 1, img_rows, img_cols)
        input_shape = (1, img_rows, img_cols)
    else:
        X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
        X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
        input_shape = (img_rows, img_cols, 1)
    #more reshaping
    X_train = X_train.astype('float32')
    X_test = X_test.astype('float32')
    X_train /= 255
    X_test /= 255
    print('X_train shape:', X_train.shape) #X_train shape: (60000, 28, 28, 1)
```

X_train shape: (60000, 28, 28, 1)

```
In [7]: num_category = 10
        y_train = keras.utils.to_categorical(y_train, num_category)
        y_test = keras.utils.to_categorical(y_test, num_category)
```

```
In [8]: ##model building
        model = Sequential()

        model.add(Conv2D(32, kernel_size=(3, 3),
                           activation='relu',
                           input_shape=input_shape))

        model.add(Conv2D(64, (3, 3), activation='relu'))

        model.add(MaxPooling2D(pool_size=(2, 2)))
```

```

model.add(Flatten())

model.add(Dense(128, activation='relu'))

model.add(Dropout(0.5))

#output a softmax to squash the matrix into output probabilities
model.add(Dense(num_category, activation='softmax'))

```

```

In [9]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])

```

```

In [10]: batch_size = 128
         num_epoch = 10
         #model training
         model_log = model.fit(X_train, y_train,
                             batch_size=batch_size,
                             epochs=num_epoch,
                             verbose=1,
                             validation_data=(X_test, y_test))

```

```

Epoch 1/10
469/469 [=====] - 79s 164ms/step - loss: 2.2910 - accuracy: 0.1269 - val_loss: 2.2671 - val_accuracy: 0.3400
Epoch 2/10
469/469 [=====] - 77s 163ms/step - loss: 2.2537 - accuracy: 0.2302 - val_loss: 2.2222 - val_accuracy: 0.4923
Epoch 3/10
469/469 [=====] - 76s 163ms/step - loss: 2.2063 - accuracy: 0.3176 - val_loss: 2.1615 - val_accuracy: 0.5636
Epoch 4/10
469/469 [=====] - 78s 166ms/step - loss: 2.1407 - accuracy: 0.3851 - val_loss: 2.0777 - val_accuracy: 0.6009
Epoch 5/10
469/469 [=====] - 73s 157ms/step - loss: 2.0527 - accuracy: 0.4445 - val_loss: 1.9640 - val_accuracy: 0.6423
Epoch 6/10
469/469 [=====] - 73s 156ms/step - loss: 1.9352 - accuracy: 0.4981 - val_loss: 1.8151 - val_accuracy: 0.6915
Epoch 7/10
469/469 [=====] - 76s 161ms/step - loss: 1.7908 - accuracy: 0.5451 - val_loss: 1.6349 - val_accuracy: 0.7268
Epoch 8/10
469/469 [=====] - 76s 162ms/step - loss: 1.6287 - accuracy: 0.5813 - val_loss: 1.4365 - val_accuracy: 0.7588
Epoch 9/10
469/469 [=====] - 74s 157ms/step - loss: 1.4634 - accuracy: 0.6148 - val_loss: 1.2443 - val_accuracy: 0.7811
Epoch 10/10
469/469 [=====] - 72s 154ms/step - loss: 1.3159 - accuracy: 0.6406 - val_loss: 1.0816 - val_accuracy: 0.7994

```

```

In [11]: score = model.evaluate(X_test, y_test, verbose=0)
         print('Test loss:', score[0])
         print('Test accuracy:', score[1])

```

```

Test loss: 1.0816175937652588
Test accuracy: 0.7993999719619751

```

```

In [12]: import os
         # plotting the metrics

```

```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js
plt.subplot(2,1,1)

```

```
plt.plot(model_log.history['accuracy'])
plt.plot(model_log.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')
plt.subplot(2,1,2)
plt.plot(model_log.history['loss'])
plt.plot(model_log.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')
plt.tight_layout()
fig
```

Out[12]:

